

France Unemployment predictions

Importing dataset:

```
rm(list=ls())
data_france=read.csv("France_Unemployment.csv")
head(data_france)
```

```
##      TIME Unemployment_Rate Unemployment    CPI Total_new_Job_Vacancies
## 1 1989-01          0.098      2554100 64.215          101700
## 2 1989-02          0.097      2541300 64.398          100000
## 3 1989-03          0.097      2533900 64.582           99500
## 4 1989-04          0.097      2547800 64.986          106100
## 5 1989-05          0.097      2537300 65.244          105300
## 6 1989-06          0.096      2530900 65.317          110600
##      Domestic_Producer_Prices_Index Unemployed_Male Unemployed_Female
## 1                                82.4663          1047000          1365000
## 2                                82.4663          1037000          1360000
## 3                                82.5494          1033000          1360000
## 4                                83.0482          1030000          1359000
## 5                                83.1313          1029000          1357000
## 6                                82.6325          1028000          1355000
```

```
tail(data_france)
```

```
##      TIME Unemployment_Rate Unemployment    CPI Total_new_Job_Vacancies
## 405 2022-09          0.071      2905800 112.74          323400
## 406 2022-10          0.072      2875100 113.90          324100
## 407 2022-11          0.071      2810400 114.26          341800
## 408 2022-12          0.071      2816600 114.16          324600
## 409 2023-01          0.071      2808500 114.60          337100
## 410 2023-02          0.070      2780800 115.78          342300
##      Domestic_Producer_Prices_Index Unemployed_Male Unemployed_Female
## 405                                144.5          1154000          1019000
## 406                                144.2          1165000          1039000
## 407                                145.2          1154000          1033000
## 408                                146.9          1157000          1033000
## 409                                150.5          1154000          1020000
## 410                                149.2          1147000          1005000
```

Feature Scaling - Creating a new variable:

```
data_france$male_to_female_unemp=round((data_france$Unemployed_Male/data_france$Unemployed_Female),4)
print(data_france$male_to_female_unemp[10])
```

```
## [1] 0.7654
```

This is to incorporate the factor - whether female are getting more unemployed or not compared to males over the years - as an external variable for overall unemployment rate.

Checking Multicollinearity using VIFs:

```
suppressWarnings(library(regclass))
```

```
## Loading required package: bestglm
```

```
## Loading required package: leaps
```

```
## Loading required package: VGAM
```

```
## Loading required package: stats4
```

```
## Loading required package: splines
```

```
## Loading required package: rpart
```

```
## Loading required package: randomForest
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
## Important regclass change from 1.3:  
## All functions that had a . in the name now have an _  
## all.correlations -> all_correlations, cor.demo -> cor_demo, etc.
```

```
VIF(lm(formula = Unemployment_Rate ~ Unemployment+CPI+Total_new_Job_Vacancies+Domestic_Producer_Pr  
ices_Index+male_to_female_unemp, data = data_france))
```

```
##                Unemployment                CPI  
##                2.359338                22.296290  
##    Total_new_Job_Vacancies Domestic_Producer_Prices_Index  
##                3.574387                6.225157  
##          male_to_female_unemp  
##                7.026014
```

Removing CPI:

```
VIF(lm(formula = Unemployment_Rate ~ Unemployment+Total_new_Job_Vacancies+Domestic_Producer_Prices  
_Index+male_to_female_unemp, data = data_france))
```

```
## Unemployment Total_new_Job_Vacancies
## 1.732834 2.039151
## Domestic_Producer_Prices_Index male_to_female_unemp
## 2.880349 3.636023
```

This is the final set of variables free from multicollinearity.

Train-Test split of the dataset - last 6 months of the data would be taken into testing part:

```
df_france_train1=data_france[1:(nrow(data_france)-6),]
df_france_test1=data_france[(nrow(data_france)-5):nrow(data_france),]
head(df_france_train1)
```

```
## TIME Unemployment_Rate Unemployment CPI Total_new_Job_Vacancies
## 1 1989-01 0.098 2554100 64.215 101700
## 2 1989-02 0.097 2541300 64.398 100000
## 3 1989-03 0.097 2533900 64.582 99500
## 4 1989-04 0.097 2547800 64.986 106100
## 5 1989-05 0.097 2537300 65.244 105300
## 6 1989-06 0.096 2530900 65.317 110600
## Domestic_Producer_Prices_Index Unemployed_Male Unemployed_Female
## 1 82.4663 1047000 1365000
## 2 82.4663 1037000 1360000
## 3 82.5494 1033000 1360000
## 4 83.0482 1030000 1359000
## 5 83.1313 1029000 1357000
## 6 82.6325 1028000 1355000
## male_to_female_unemp
## 1 0.7670
## 2 0.7625
## 3 0.7596
## 4 0.7579
## 5 0.7583
## 6 0.7587
```

- The model would be trained on the train dataset.
- And the performance of the fitted model would be checked on the test dataset.
- If this performs fairly well, this model would be considered to get the future forecasts.

Time series plot:

```
suppressWarnings(library(fpp2))
```

```
## Registered S3 method overwritten by 'quantmod':
## method from
## as.zoo.data.frame zoo
```

```
## — Attaching packages — fpp2 2.5 —
```

```
## ✓ ggplot2 3.3.6 ✓ fma 2.4
## ✓ forecast 8.18 ✓ expsmooth 2.3
```

```
## — Conflicts ————— fpp2_conflicts —
## ✖ ggplot2::margin() masks randomForest::margin()
```

```
suppressWarnings(library(urca))
df.ts=ts(data_france$Unemployment_Rate, frequency = 12, start = c(1989,1))
plot(df.ts,xlab="Years",ylab="Unemployment Rates(Proportions)")
title(main="Time series plot of unemployment rate in France")
```

Time series plot of unemployment rate in France



Testing stationarity:

```
df_france_train1[, "Unemployment_Rate"] %>%
  ur.kpss() %>%
  summary()
```

```
##
## #####
## # KPSS Unit Root Test #
## #####
##
## Test is of type: mu with 5 lags.
##
## Value of test-statistic is: 2.3638
##
## Critical value for a significance level of:
##          10pct  5pct 2.5pct  1pct
## critical values 0.347 0.463  0.574 0.739
```

This series is non-stationary - 1st order differencing would be necessary.

Testing stationarity after 1st order differencing:

```
diff(df_france_train1[, "Unemployment_Rate"]) %>%
  ur.kpss() %>%
  summary()
```


```
##
## #####
## # KPSS Unit Root Test #
## #####
##
## Test is of type: mu with 5 lags.
##
## Value of test-statistic is: 0.2404
##
## Critical value for a significance level of:
##          10pct  5pct 2.5pct  1pct
## critical values 0.347 0.463  0.574 0.739
```

1st order differences are stationary.

ACF plot:

```
par(mfrow=c(1,1))
acf(diff(df_france_train1$Unemployment_Rate), lag.max = 20, main = "ACF plot")
```

ACF plot




p can be taken as 0/1/2 based on the no. of significant lags.

PACF plot:

```
par(mfrow=c(1,1))
pacf(diff(df_france_train1$Unemployment_Rate), lag.max = 10, main = "PACF plot")
```

PACF plot



q can be 0/1/2, based on the no. of significant lags.

Fitting ARIMAX model ignoring the variables that were eliminated due to high VIF:

Starting with the value of p & q as 2 and with the rest of the regressors:

```
est_train=arima(df_france_train1$Unemployment_Rate, order=c(2,1,2), xreg = as.matrix(df_france_train1[,c(3,5,6,9)]), method = "ML")
summary(est_train)
```

```
##
## Call:
## arima(x = df_france_train1$Unemployment_Rate, order = c(2, 1, 2), xreg = as.matrix(df_france_train1[,
##      c(3, 5, 6, 9)]), method = "ML")
##
## Coefficients:
```

```
## Warning in sqrt(diag(x$var.coef)): NaNs produced
```

```
##          ar1          ar2          ma1          ma2  Unemployment  Total_new_Job_Vacancies
##      0.8795 -0.6445 -0.2222  0.3565              0              0
## s.e.  0.0959  0.0559  0.1065  0.0705              NaN              NaN
##      Domestic_Producer_Prices_Index  male_to_female_unemp
##                                  0e+00              -0.0093
## s.e.                                  1e-04              0.0039
##
## sigma^2 estimated as 1.049e-06:  log likelihood = 2201.9,  aic = -4385.81
##
## Training set error measures:
##              ME              RMSE              MAE              MPE              MAPE
## Training set -4.424217e-05  0.001022971  0.0006904664 -0.05289185  0.7428873
##              MASE              ACF1
## Training set  0.875025 -0.008829127
```

Test of significance of individual coefficients:

```
suppressWarnings(library(lmtest))
```

```
## Loading required package: zoo
```

```
##
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':
##
##      as.Date, as.Date.numeric
```

```
##
## Attaching package: 'lmtest'
```

```
## The following object is masked from 'package:VGAM':
##
##      lrtest
```

```
coeftest(est_train)
```

```
## Warning in sqrt(diag(se)): NaNs produced
```



```
##
## z test of coefficients:
##
##               Estimate Std. Error  z value  Pr(>|z|)
## ar1             8.7946e-01  9.5888e-02   9.1717 < 2.2e-16 ***
## ar2            -6.4447e-01  5.5925e-02 -11.5239 < 2.2e-16 ***
## ma1            -2.2216e-01  1.0651e-01  -2.0858  0.03700 *
## ma2             3.5654e-01  7.0496e-02   5.0576 4.245e-07 ***
## Unemployment     1.1161e-08         NaN      NaN      NaN
## Total_new_Job_Vacancies -3.3634e-09         NaN      NaN      NaN
## Domestic_Producer_Prices_Index -2.8934e-06  9.9732e-05  -0.0290  0.97686
## male_to_female_unemp    -9.3459e-03  3.9176e-03  -2.3856  0.01705 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

We need to remove the variables producing NaNs & the insignificant variables.

After doing that, the summary & test of significances of the final model would look like:

```
est_1=arima(df_france_train1$Unemployment_Rate, order=c(1,1,1), method = "ML")
summary(est_1)
```

```
##
## Call:
## arima(x = df_france_train1$Unemployment_Rate, order = c(1, 1, 1), method = "ML")
##
## Coefficients:
##          ar1      ma1
##      0.2906  0.1905
## s.e.  0.0986  0.0994
##
## sigma^2 estimated as 1.335e-06:  log likelihood = 2153.59,  aic = -4301.19
##
## Training set error measures:
##              ME          RMSE          MAE          MPE          MAPE
## Training set -3.973705e-05  0.001154181  0.0007430032 -0.05063419  0.8069906
##
##              MASE          ACF1
## Training set  0.9416046  0.002304672
```

Test of significance of coefficients:

```
suppressWarnings(library(lmtest))
coeftest(est_1)
```

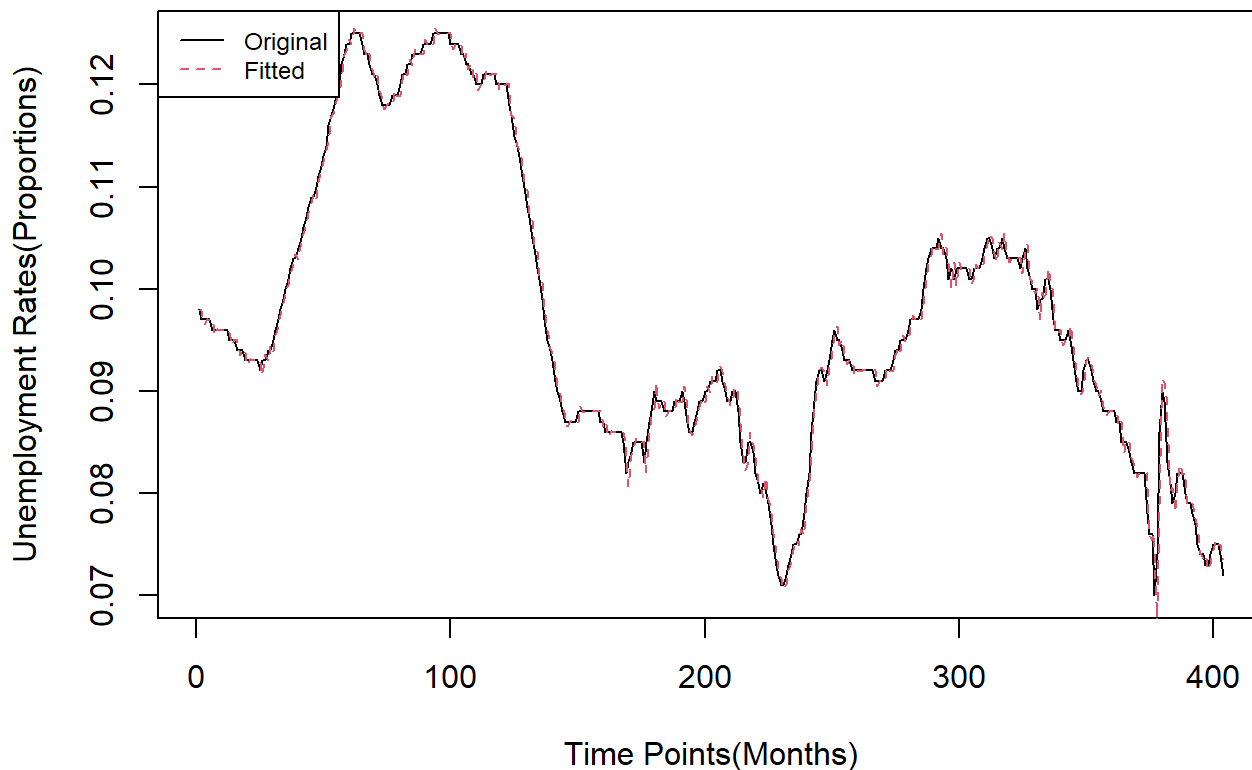
```
##
## z test of coefficients:
##
##      Estimate Std. Error z value Pr(>|z|)
## ar1 0.290572    0.098581  2.9475 0.003203 **
## ma1 0.190456    0.099390  1.9163 0.055332 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Thus only those parameters are kept as final, which are significant in prediction of the target variable.

Plot of Fitted vs Original values for train dataset:

```
res=residuals(est_1)
data_fit=df_france_train1$Unemployment_Rate-res
ts.plot(df_france_train1$Unemployment_Rate, type="l", xlab="Time Points(Months)", ylab="Unemployment Rates(Proportions)", main="Fitted vs original for train dataset")
points(data_fit, type="l", col=2, lty=2)
legend("topleft",c("Original","Fitted"), col=c(1,2), lty=c(1,2), cex=0.75)
```

Fitted vs original for train dataset



Predictions of unemployment rates for the test dataset using above fitted model:

```
test_pred=predict(est_1, n.ahead=6, se.fit=FALSE, method="ML")
```

Predicted values:

```
print(as.vector(test_pred))
```

```
## [1] 0.07113007 0.07087730 0.07080385 0.07078251 0.07077630 0.07077450
```

Original values:

```
print(df_france_test1$Unemployment_Rate)
```

```
## [1] 0.071 0.072 0.071 0.071 0.071 0.070
```

Performance on test dataset:

MAPE (in %):

```
(1/length(df_france_test1$Unemployment_Rate))*(sum(abs(df_france_test1$Unemployment_Rate-as.vector(test_pred))/abs(df_france_test1$Unemployment_Rate)))*100
```

```
## [1] 0.6244345
```

RMSE:

```
sqrt(mean((df_france_test1$Unemployment_Rate-as.vector(test_pred))^2))
```

```
## [1] 0.000579231
```

Thus, the fitted model is working well, more or less, for future dataset.

Now going with the same approach with the actual dataset for getting the future forecast of March,23:

Checking stationarity:

```
data_france[, "Unemployment_Rate"] %>%  
  ur.kpss() %>%  
  summary()
```

```
##
## #####
## # KPSS Unit Root Test #
## #####
##
## Test is of type: mu with 5 lags.
##
## Value of test-statistic is: 2.5039
##
## Critical value for a significance level of:
##          10pct  5pct 2.5pct  1pct
## critical values 0.347 0.463  0.574 0.739
```


```
diff(data_france[, "Unemployment_Rate"]) %>%
  ur.kpss() %>%
  summary()
```

```
##
## #####
## # KPSS Unit Root Test #
## #####
##
## Test is of type: mu with 5 lags.
##
## Value of test-statistic is: 0.2521
##
## Critical value for a significance level of:
##          10pct  5pct 2.5pct  1pct
## critical values 0.347 0.463  0.574 0.739
```

ACF plot:

```
par(mfrow=c(1,1))
acf(diff(data_france$Unemployment_Rate), lag.max = 20, main = "ACF plot")
```

ACF plot




p can be taken as 0/1/2, based on the no. of significant lags.

PACF plot:

```
par(mfrow=c(1,1))
pacf(diff(data_france$Unemployment_Rate), lag.max = 20, main = "PACF plot")
```

PACF plot



```
#q=0/1
```

q can be taken as 0/1, based on the no. of significant lags.

Fitting the model that we tested before - on the actual data:

```
est_actual=arima(data_france$Unemployment_Rate, order=c(1,1,1), method = "ML")
summary(est_actual)
```

```
##
## Call:
## arima(x = data_france$Unemployment_Rate, order = c(1, 1, 1), method = "ML")
##
## Coefficients:
##          ar1      ma1
##      0.2849  0.1913
## s.e.  0.0994  0.1004
##
## sigma^2 estimated as 1.329e-06:  log likelihood = 2186.63,  aic = -4367.27
##
## Training set error measures:
##              ME          RMSE          MAE          MPE          MAPE
## Training set -4.15456e-05 0.001151466 0.000743623 -0.05339375 0.8113157
##              MASE          ACF1
## Training set 0.9445397 0.001563919
```


Test of significance of individual coefficients:

```
suppressWarnings(library(lmtest))
coeftest(est_actual)
```

```
##
## z test of coefficients:
##
##      Estimate Std. Error z value Pr(>|z|)
## ar1 0.284904    0.099398  2.8663 0.004153 **
## ma1 0.191338    0.100386  1.9060 0.056646 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Plot of Fitted vs Original on the actual data:

```
res=residuals(est_actual)
data_fit=data_france$Unemployment_Rate-res
ts.plot(data_france$Unemployment_Rate, type="l", xlab="Time Points(Months)", ylab="Unemployment Ra
tes(Proportions)", main="Fitted vs original for France")
points(data_fit, type="l", col=2, lty=2)
legend("topright",c("Original","Fitted"), col=c(1,2), lty=c(1,2), cex=0.75)
```



Obtaining prediction of Unemployment rate for April 2023:

```
#Adding actual unemployment rate of March 23:  
unemp_fr=c(data_france$Unemployment_Rate,0.069)  
est_actual=arima(unemp_fr, order=c(1,1,1), method = "ML")  
  
future_unemp_pred=predict(est_actual, n.ahead=1, se.fit=FALSE, method="ML")  
print(as.vector(future_unemp_pred))
```

```
## [1] 0.06861347
```

Upper & Lower limits (95% C.I.s):

```
upper=as.vector(future_unemp_pred)+(1.96*(sqrt(est_actual$sigma2)))  
lower=as.vector(future_unemp_pred)-(1.96*(sqrt(est_actual$sigma2)))
```

Upper limit for April 2023 forecast:

```
print(as.vector(upper))
```

```
## [1] 0.0708709
```

Lower limit for April 2023 forecast:

```
print(as.vector(lower))
```

```
## [1] 0.06635605
```

April,23 forecast - 6.8613 %

Upper & Lower limits - (6.6356 %, 7.087 %)